

TiDB Operator 实现原理解析

weekface@PingCAP



扫码关注公众号,获取 316 北京站 PPT

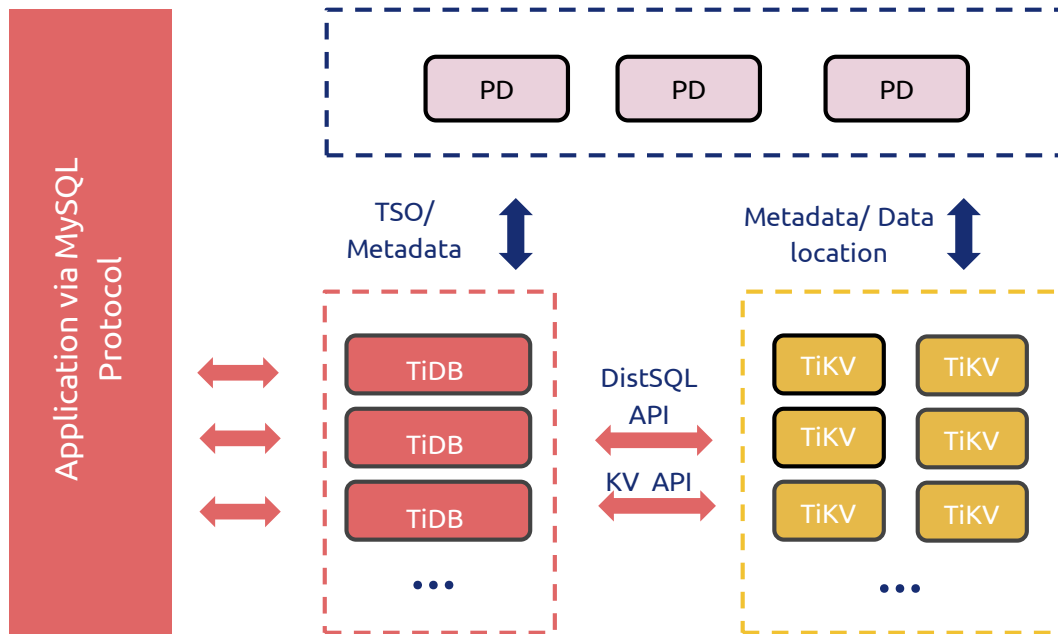
About me

weekface (常军昌), software engineer, PingCAP

Focus on Cloud TiDB

changjunchang@pingcap.com

About TiDB



Traditional Deployment

PD1

```
pd-server --initial=pd1 --data-dir=...
```

Traditional Deployment

PD1

`pd-server --initial=pd1 --data-dir=...`

PD2

`pd-server --join=pd1 --data-dir=...`

Traditional Deployment

PD1

`pd-server --initial=pd1 --data-dir=...`

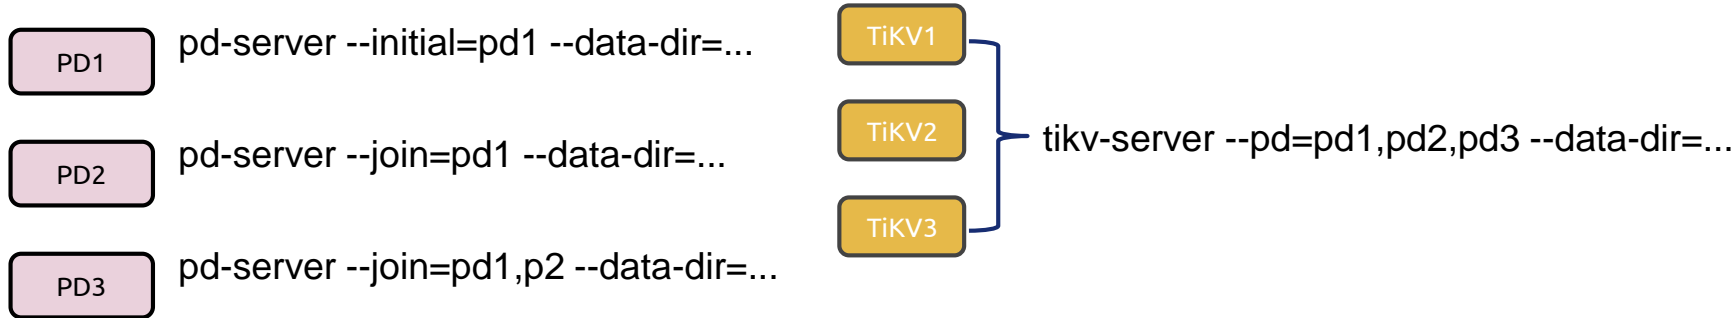
PD2

`pd-server --join=pd1 --data-dir=...`

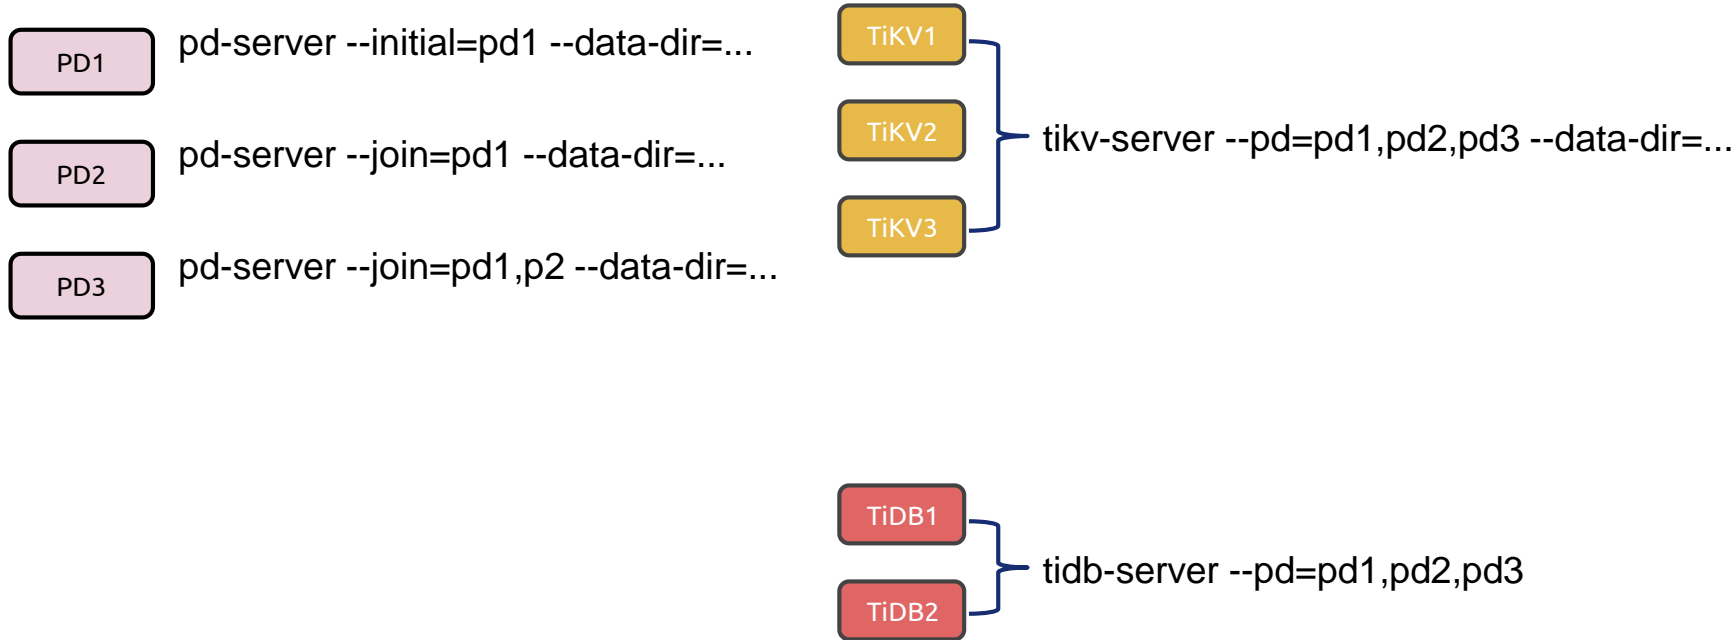
PD3

`pd-server --join=pd1,p2 --data-dir=...`

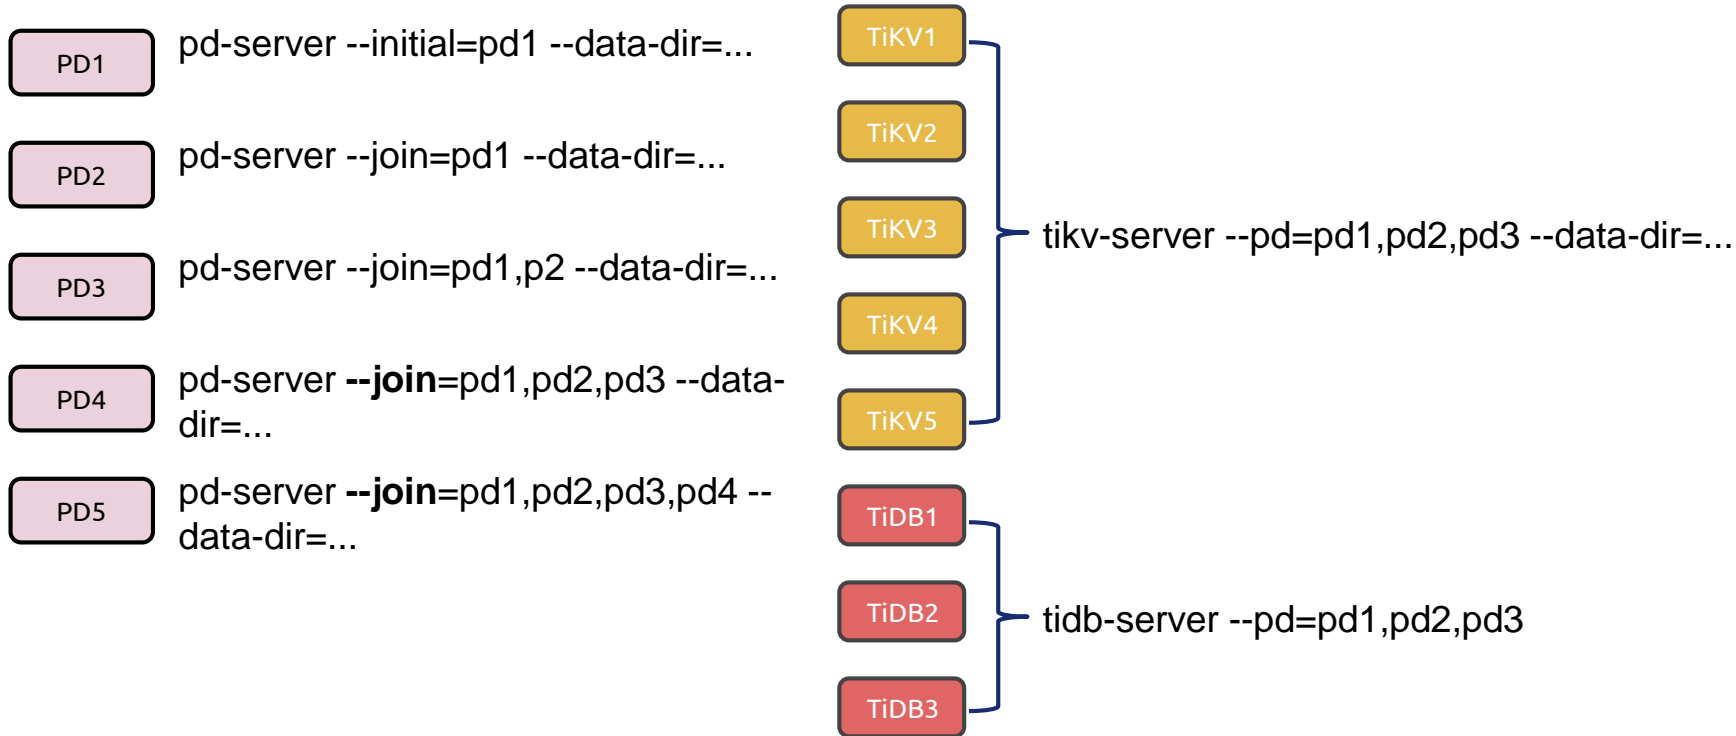
Traditional Deployment



Traditional Deployment



Traditional Deployment



Summary

Operation	PD	TiKV	TiDB
Persistent storage	Yes	Yes	-
Create/Upgrade Order	1	2	3
Startup args	Different	Same	Same
Before Stop/Offline	Resign Leader/ Delete Member	Evict Leader/ Delete Store	Resign DDL Owner
Failover			

What's TiDB Operator?

TiDB operator creates and manages TiDB clusters running in Kubernetes



Standing on the shoulders of giants

Based on Raw **Pod**, we can do all the things. But we don't want to recreate a wheel

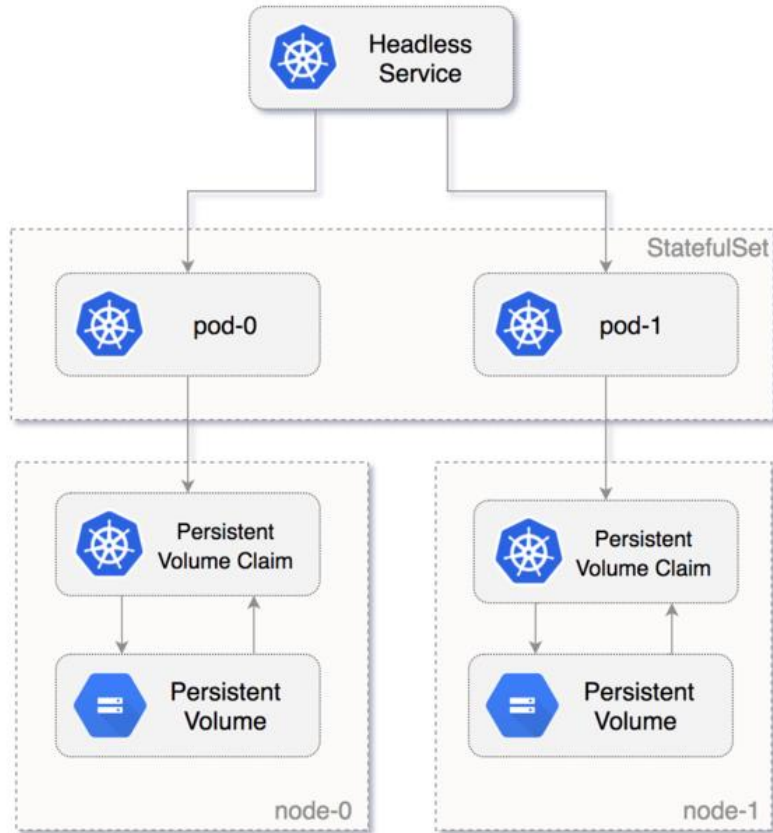
Standing on the shoulders of giants

ReplicaSet, **Deployment** and other workloads are used to manage **stateless** applications

Standing on the shoulders of giants

StatefulSet is the workload API object used to manage **stateful** applications.

- Stable, unique network identifiers.
- Stable, persistent storage.
- Ordered, graceful deployment and scaling.
- Ordered, automated rolling updates.




Based on StatefulSet - Persistent storage

PersistentVolumes associated with the Pods' PersistentVolumeClaims


Based on StatefulSet - Create/Upgrade Order

Create and manage three StatefulSet objects: PD StatefulSet, TiKV StatefulSet, TiDB StatefulSet


Based on StatefulSet - PD Startup Args



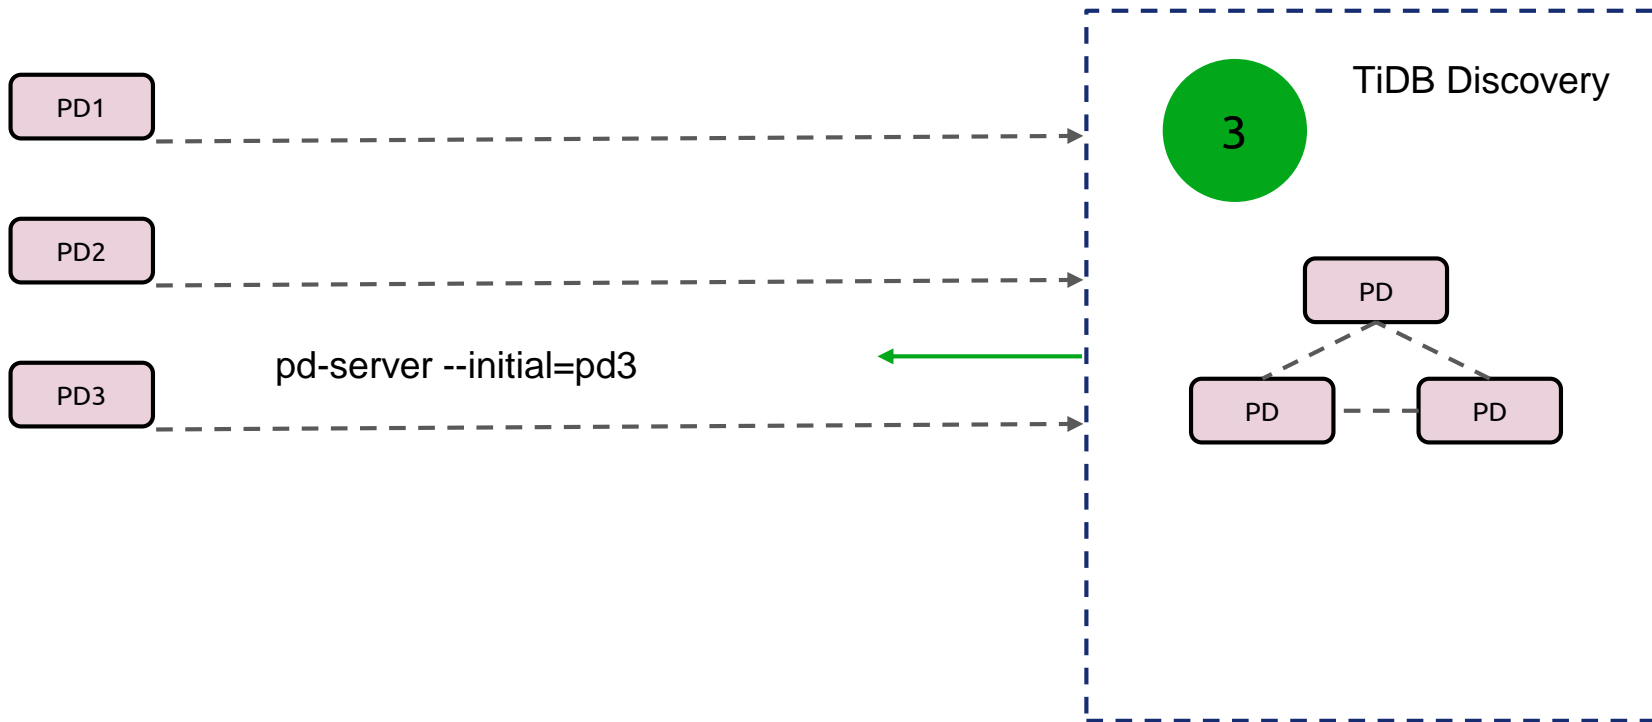
Based on StatefulSet - PD Startup Args



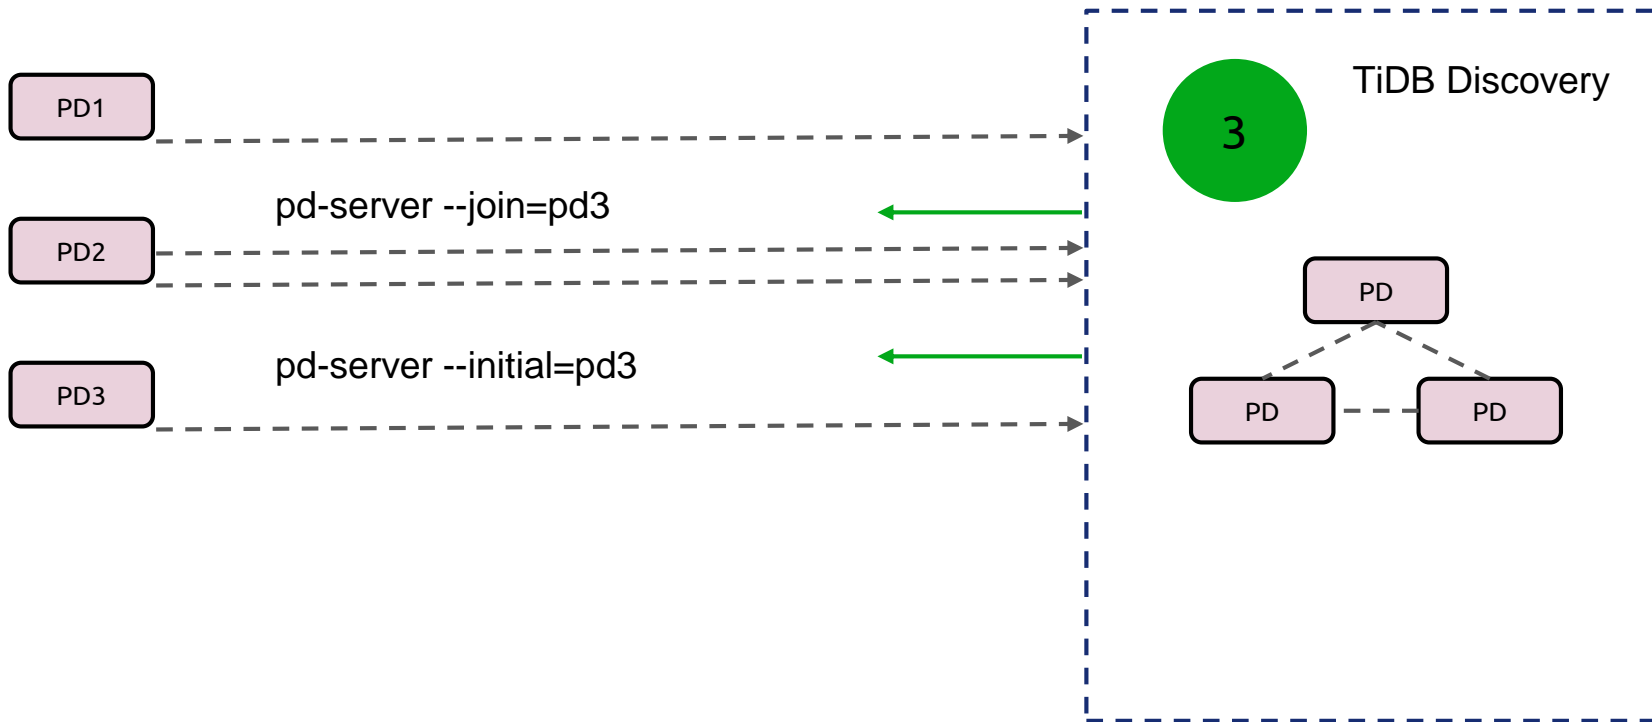
Based on StatefulSet - PD Startup Args



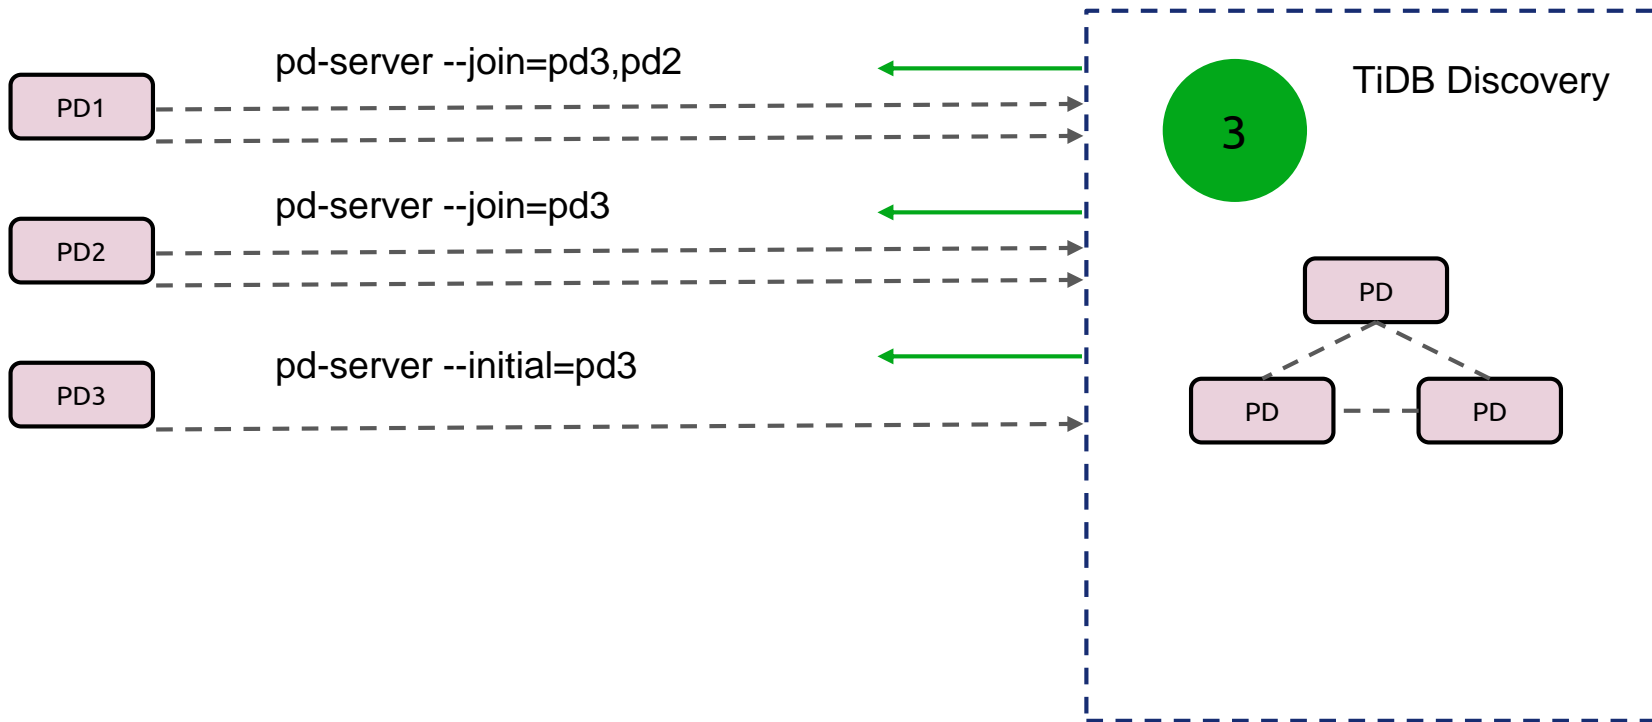
Based on StatefulSet - PD Startup Args




Based on StatefulSet - PD Startup Args



Based on StatefulSet - PD Startup Args



Based on StatefulSet - PD Startup Args



Based on StatefulSet - Before Stop/Offline

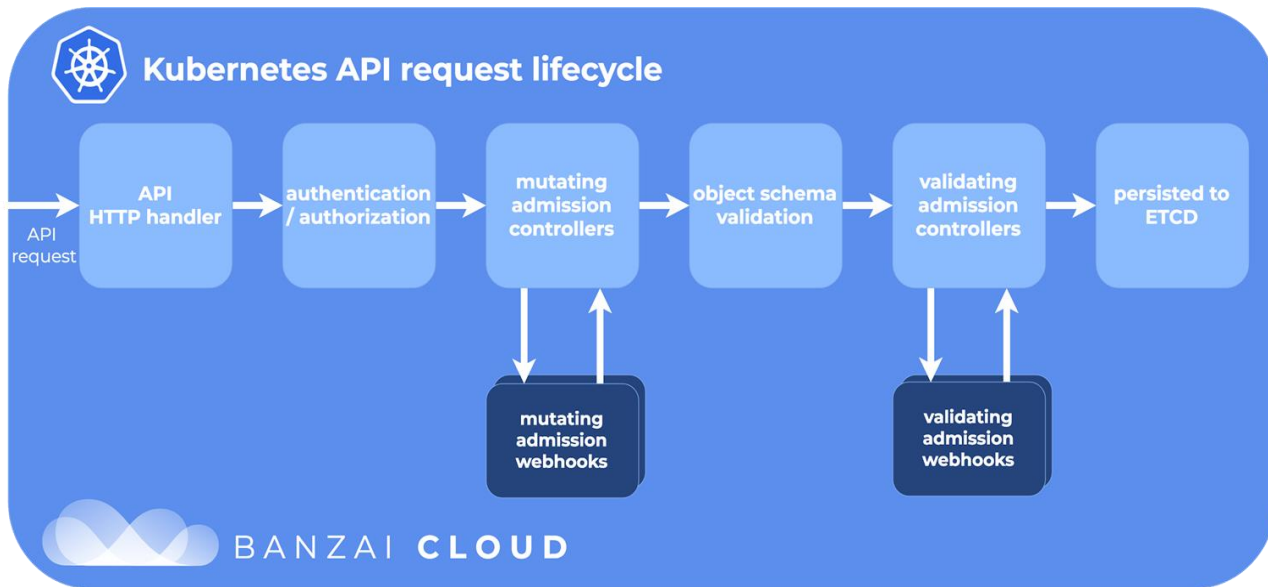
We should do something before the Pod terminated.

- K8s PreStop probe
- StatefulSet [partition](#)
- Validating admission webhooks

Based on StatefulSet - Before Stop/Offline

We should do something before the Pod terminated.

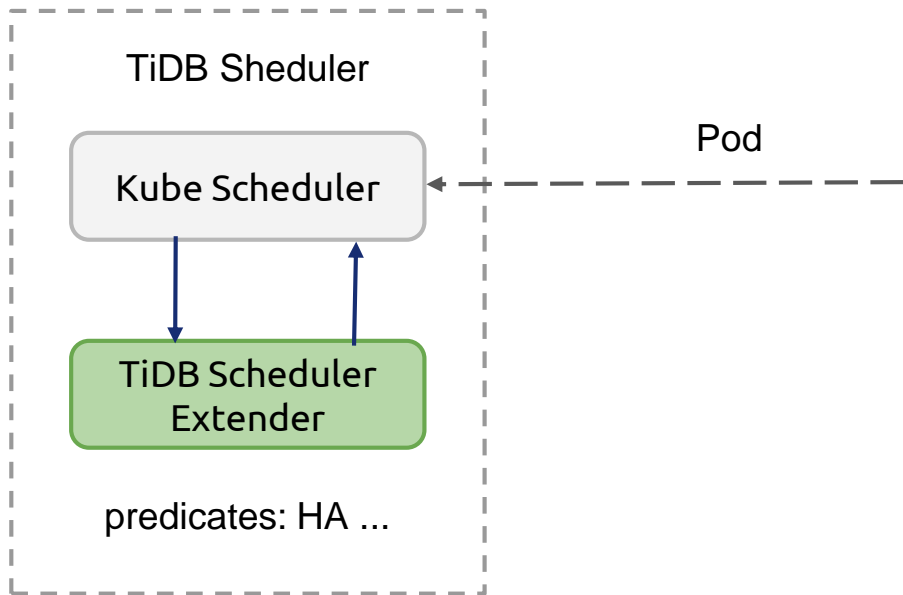
- K8s PreStop probe
- StatefulSet [partition](#)
- Validating admission webhooks



TiDB Scheduler


For data safety, we don't permit two TiKV instances scheduled into the same node

```
apiVersion: apps/v1
kind: StatefulSet
...
spec:
  template:
    spec:
      schedulerName: tidb-scheduler
    containers:
  ...
```




Failover

StatefulSet don't failover automatically, we need this feature for TiDB cluster.



Failover

StatefulSet don't failover automatically, we need this feature for TiDB cluster.



So we need a new API Object(CRD): TidbCluster

apiVersion: pingcap.com/v1alpha1

kind: **TidbCluster**

metadata:

name: demo

spec:

pd:

image: pingcap/pd:v2.1.3

replicas: 3

...

tikv:

image: pingcap/tikv:v2.1.3

replicas: 5

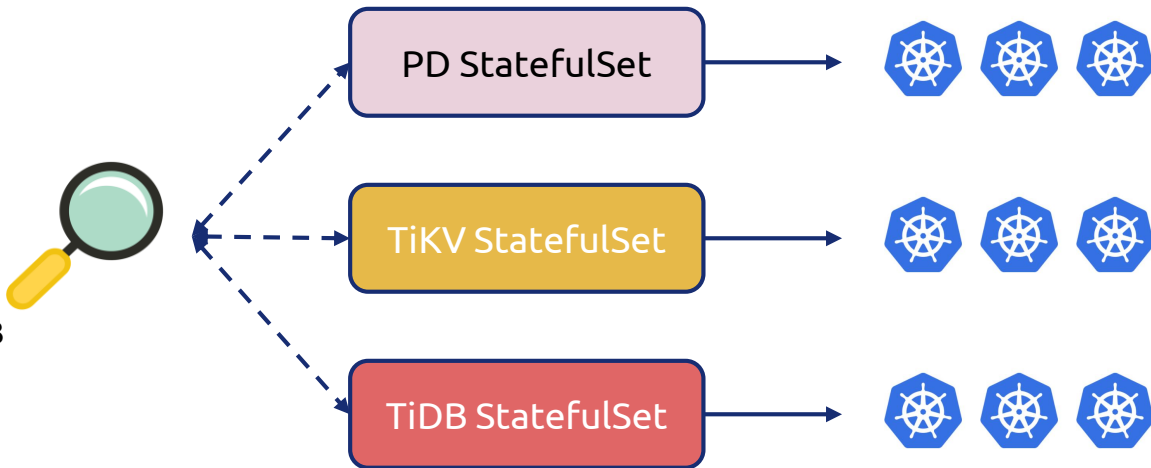
...

tidb:

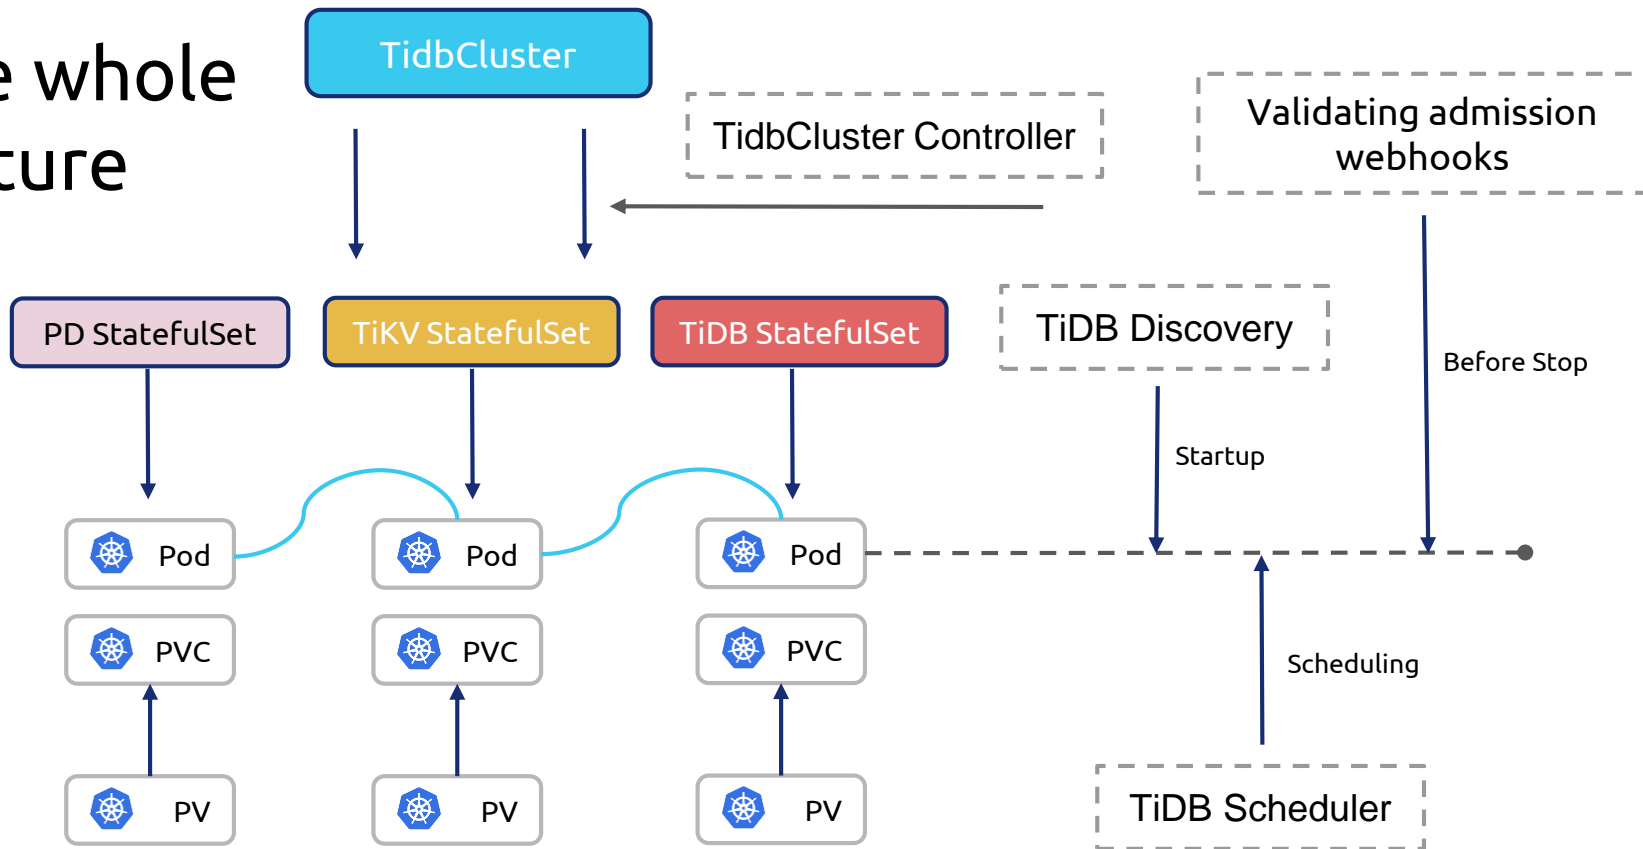
image: pingcap/tidb:v2.1.3

replicas: 2

...



The whole picture



Thank you!

<https://github.com/pingcap/tidb-operator>



扫码关注公众号,获取 316 北京站 PPT