

# 使用Kubernetes运行MXNet和AutoTVM

苏磊  
TuSimple



扫码关注公众号,获取 316 北京站 PPT

# 自我介绍

## • 工作背景:

- 2008~2011: Platform Computing (Acquired by IBM)
- 2011~2012: 腾讯云平台
- 2012~2017: IBM Spectrum Computing
- 2017~至今: 图森

## • 技能领域:

- 10+年分布式计算研发, 将分布式计算应用于高性能计算、云计算、大数据和车载系统等领域
- 资源管理、资源调度、作业调度
- 高性能车载中间件、深度学习优化

(19) **United States**

(12) **Patent Application Publication**  
**Chin et al.**

(10) **Pub. No.: US 2016/0306662 A1**

(43) **Pub. Date: Oct. 20, 2016**

(54) **MULTI-DIMENSION SCHEDULING AMONG  
MULTIPLE CONSUMERS**

**Publication Classification**

(71) Applicant: **International Business Machines  
Corporation**, Armonk, NY (US)

(51) **Int. Cl.**  
**G06F 9/50** (2006.01)

(52) **U.S. Cl.**  
CPC ..... **G06F 9/50** (2013.01)

(72) Inventors: **Alicia E. Chin**, Markham (CA);  
**Michael Feiman**, Richmond Hill (CA);  
**Zhenhua Hu**, Toronto (CA); **Jason T.  
S. Lam**, Markham (CA); **Zhimin Lin**,  
Scarborough (CA); **Lei Su**, Beijing  
(CN); **Hao Zhou**, Toronto (CA)

(57) **ABSTRACT**

Embodiments of the present invention provide systems and methods for allocating multiple resources. In one embodiment, a configured resource plan is used to construct a hierarchical tree. The system then identifies a set of unowned resources from the configured resource plan and sends the set of unowned resource to a share pool. The share pool is either a global or local pool and can be accessed by one or more consumers. In response to changes in workload demands, a set of unused resources are lent to a global or local pool.


(21) Appl. No.: **14/690,866**

(22) Filed: **Apr. 20, 2015**

100

## Why Conductor : Performance – most recent STAC Benchmark

### Visualization – Use Case 1



#### Key Points

- Conductor continues to demonstrate performance and throughput leadership driving ROI for client

#### UC1 Throughput

- 56% higher than YARN
- 57% higher than Mesos

#### UC2 Throughput

- 30% higher than YARN
- 62% higher than Mesos

#### UC3 Throughput

- 55% higher than YARN
- 88% higher than Mesos

#### UC4 –Throughput

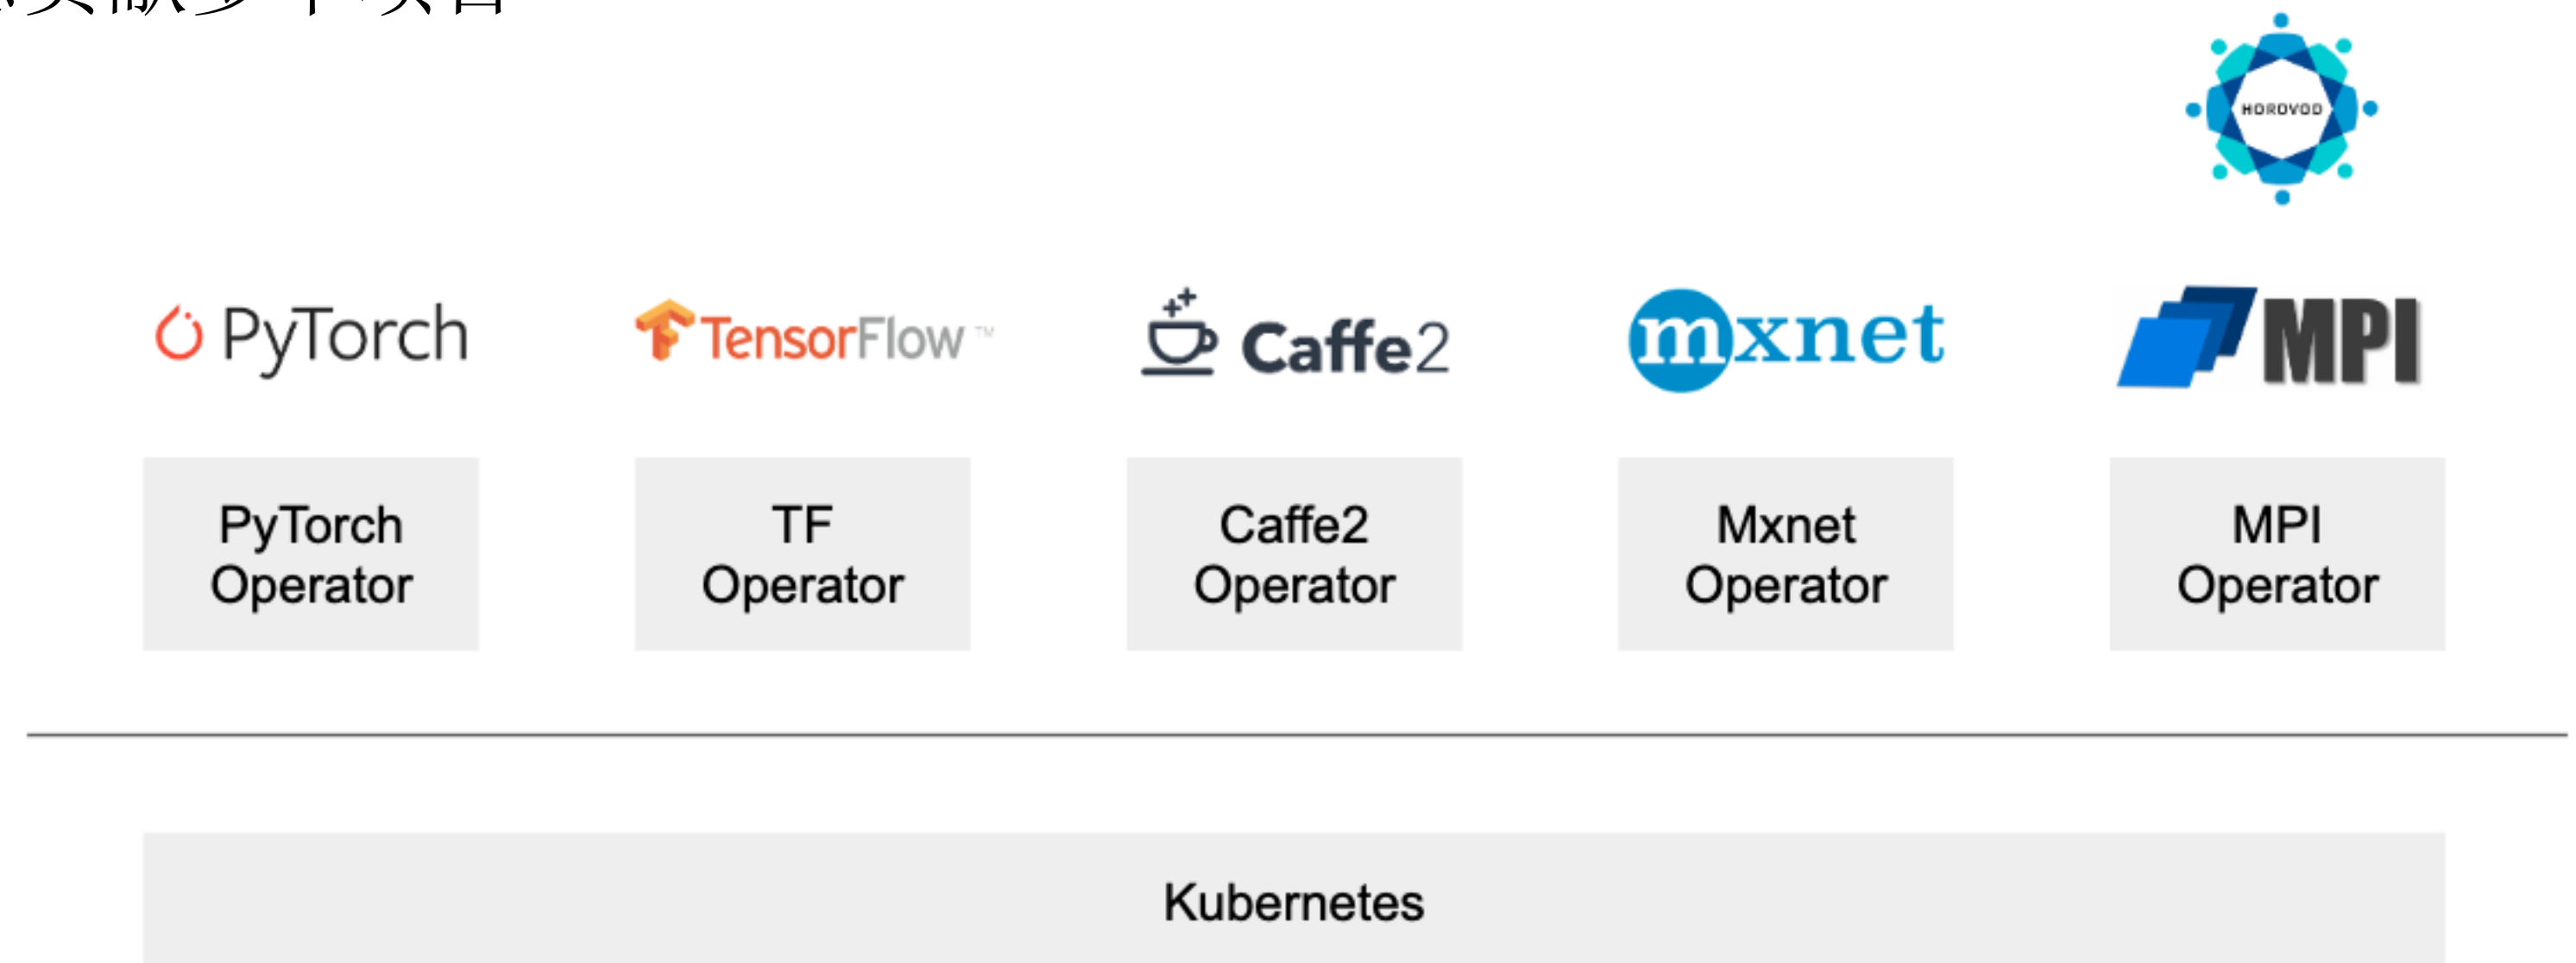
- 224% higher than YARN
- 25% higher than Mesos

- ❖ KubeFlow介绍
- ❖ MXNet介绍和分布式训练
- ❖ 使用mxnet-operator提交training任务
- ❖ TVM和AutoTVM
- ❖ 使用mxnet-operator提交tuning任务
- ❖ mxnet-operator roadmap



# 什么是Kubeflow

- [Kubeflow](#)旨在支持多种机器学习框架运行在 Kubernetes 之上，比如 Tensorflow, Pytorch, MXNet, Caffe等深度学习框架。
  - 包含operator、pipeline、超参数调优、serving 等诸多模块。
  - 通过operator提供作业生命周期管理，满足分布式训练需求，达到一键式部署训练任务的目的。
- 众多国内公司参与、为Kubeflow生态贡献多个项目
  - 图森：mxnet-operator
  - 才云：tf-operator
  - Momenta：caffe2-operator



# 什么是MXNet

- Apache MXNet是一个深度学习框架，旨在提高效率和灵活性。允许混合符号和命令式编程，以最大限度地提高效率 and 生产力。MXNet的核心是一个动态依赖调度程序，可以动态地自动并行化符号和命令操作。最重要的图优化层使符号执行更快，内存效率更高。
- 特点
  - 灵活的编程模型：支持命令式和符号式编程模型以最大化效率和性能。
  - 从云端到客户端可移植：可运行于多CPU、多GPU、集群、服务器、工作站甚至移动智能手机。
  - 多语言支持：支持七种主流编程语言，包括C++、Python、R、Scala、Julia、Matlab和JavaScript。
  - 分布式训练：支持在多CPU/GPU设备上的分布式训练，使其可充分利用云计算的规模优势。
  - 性能优化：使用一个优化的C++后端引擎并行I/O和计算，无论使用哪种语言都能达到最佳性能。
  - 云端友好：可直接与S3，HDFS和Azure兼容
- 官网地址
  - <https://mxnet.apache.org/>

# MXNet用户





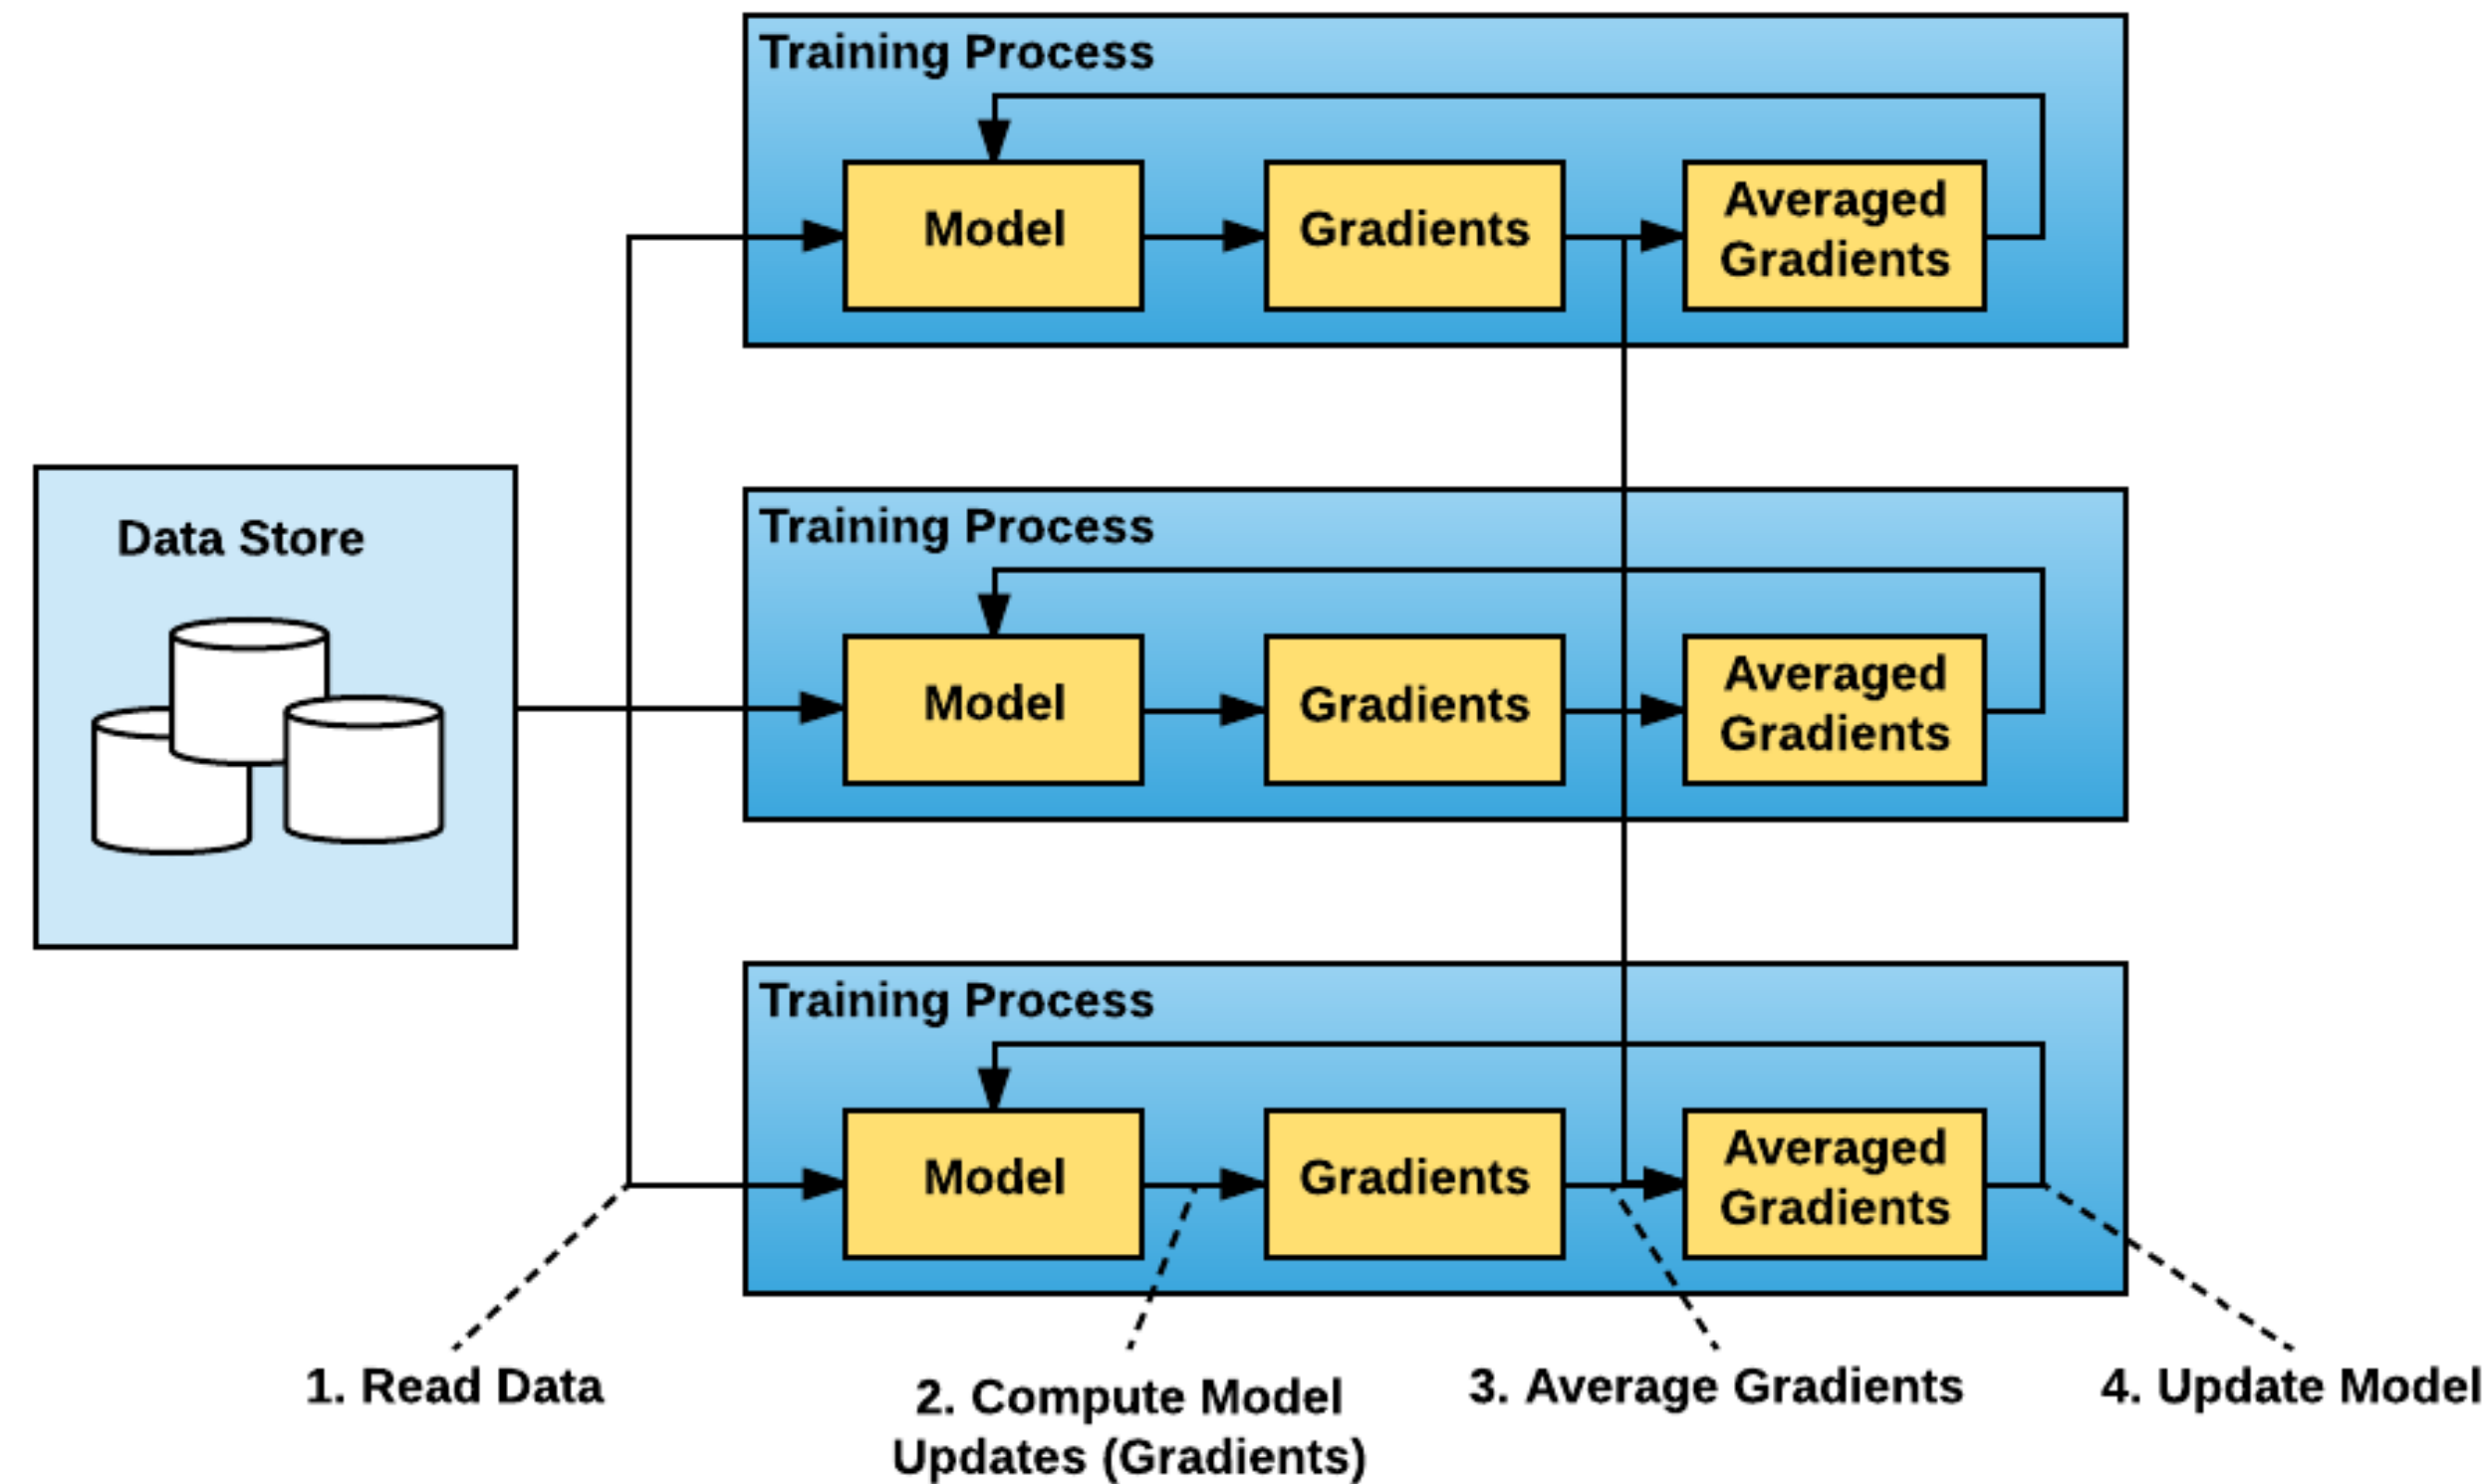
# 深度学习Pipeline





# MXNet分布式训练

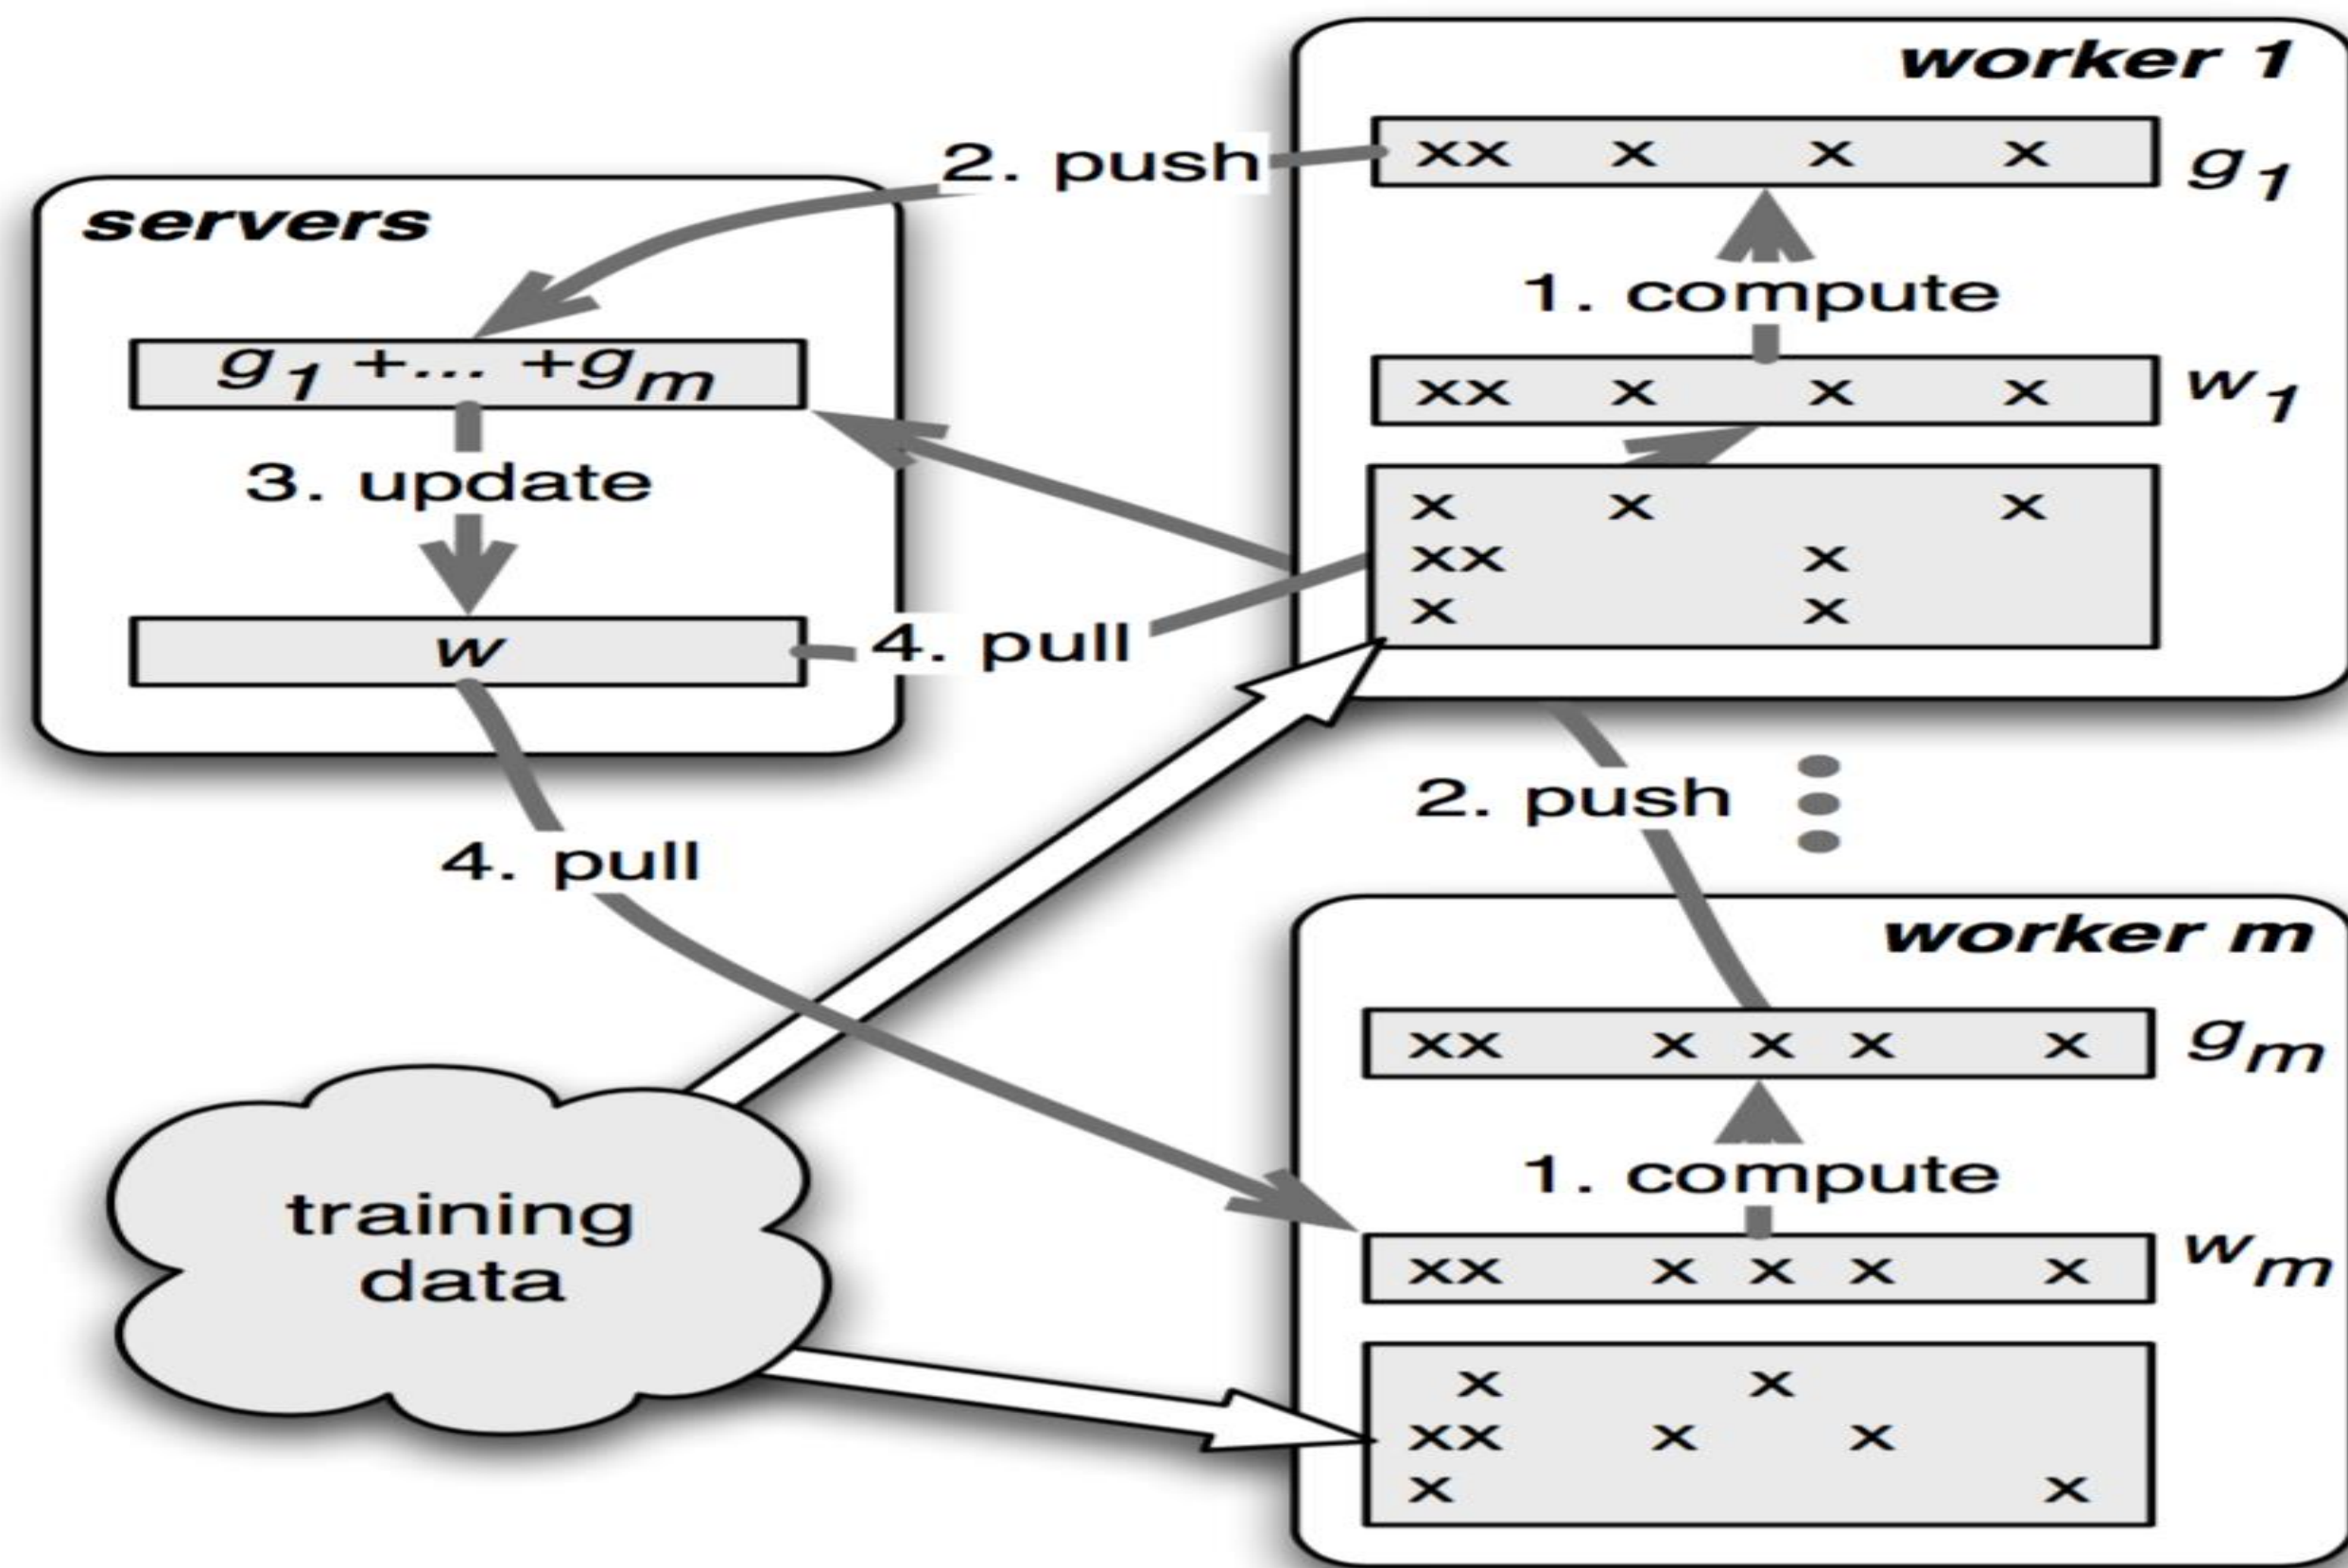
- 并行模式
  - 数据并行化：每个设备(如GPU)有一个完整的模型副本，每个设备训练一小部分数据集，共同更新共享模型，可支持单机多卡和多机多卡
  - 模型并行化：单个设备无法装下整个模型，每个设备只含有模型的一部分并且负责训练这部分模型，当前仅支持单机多卡






# MXNet分布式训练

- 通信模式
- 参数服务器




Averages All the Gradients



or

Each Averages Portion of the Gradients



# MXNet分布式训练

- 组件：
  - Server: 保存模型参数，聚合梯度更新参数。一个或多个。
  - Worker: 对训练样本进行实际训练。训练前，从参数服务器拉取参数；训练后，发送梯度到参数服务器。一个或多个。
  - Scheduler: 负责建立通信并分配和监控任务。有且仅有一个。

- 启动方式:

```
export COMMAND='python example/gluon/image_classification.py --dataset cifar10 --model vgg11  
--epochs 1 --kvstore dist_sync'
```

```
DMLC_ROLE=server DMLC_PS_ROOT_URI=127.0.0.1 DMLC_PS_ROOT_PORT=9092 DMLC_NUM_SERVER=2  
DMLC_NUM_WORKER=2 $COMMAND &
```

```
DMLC_ROLE=server DMLC_PS_ROOT_URI=127.0.0.1 DMLC_PS_ROOT_PORT=9092 DMLC_NUM_SERVER=2  
DMLC_NUM_WORKER=2 $COMMAND &
```

```
DMLC_ROLE=scheduler DMLC_PS_ROOT_URI=127.0.0.1 DMLC_PS_ROOT_PORT=9092 DMLC_NUM_SERVER=2  
DMLC_NUM_WORKER=2 $COMMAND &
```

```
DMLC_ROLE=worker DMLC_PS_ROOT_URI=127.0.0.1 DMLC_PS_ROOT_PORT=9092 DMLC_NUM_SERVER=2  
DMLC_NUM_WORKER=2 $COMMAND &
```

```
DMLC_ROLE=worker DMLC_PS_ROOT_URI=127.0.0.1 DMLC_PS_ROOT_PORT=9092 DMLC_NUM_SERVER=2  
DMLC_NUM_WORKER=2 $COMMAND
```



# 使用mxnet-operator提交training任务

```
apiVersion:
"kubeflow.org/v1alpha2"
kind: "MXJob"
metadata:
  name: "mxnet-job"
spec:
  jobMode: MXTrain
  mxReplicaSpecs:
    Scheduler:
      replicas: 1
      restartPolicy: Never
      template:
        spec:
          containers:
            - name: mxnet
              image:
mxjob/mxnet:gpu
```

```
Worker:
  replicas: 1
  restartPolicy: Never
  template:
    spec:
      containers:
        - name: mxnet
          image: mxjob/mxnet:gpu
          command: ["python"]
          args: ["/incubator-
mxnet/example/image-
classification/train_mnist.py", "--
num-epochs", "1", "--num-
layers", "2", "--kv-
store", "dist_device_sync", "--
gpus", "0"]
      resources:
        limits:
          nvidia.com/gpu: 1
```

```
Server:
  replicas: 1
  restartPolicy: Never
  template:
    spec:
      containers:
        - name: mxnet
          image:
mxjob/mxnet:gpu
```

# 使用mxnet-operator提交training任务

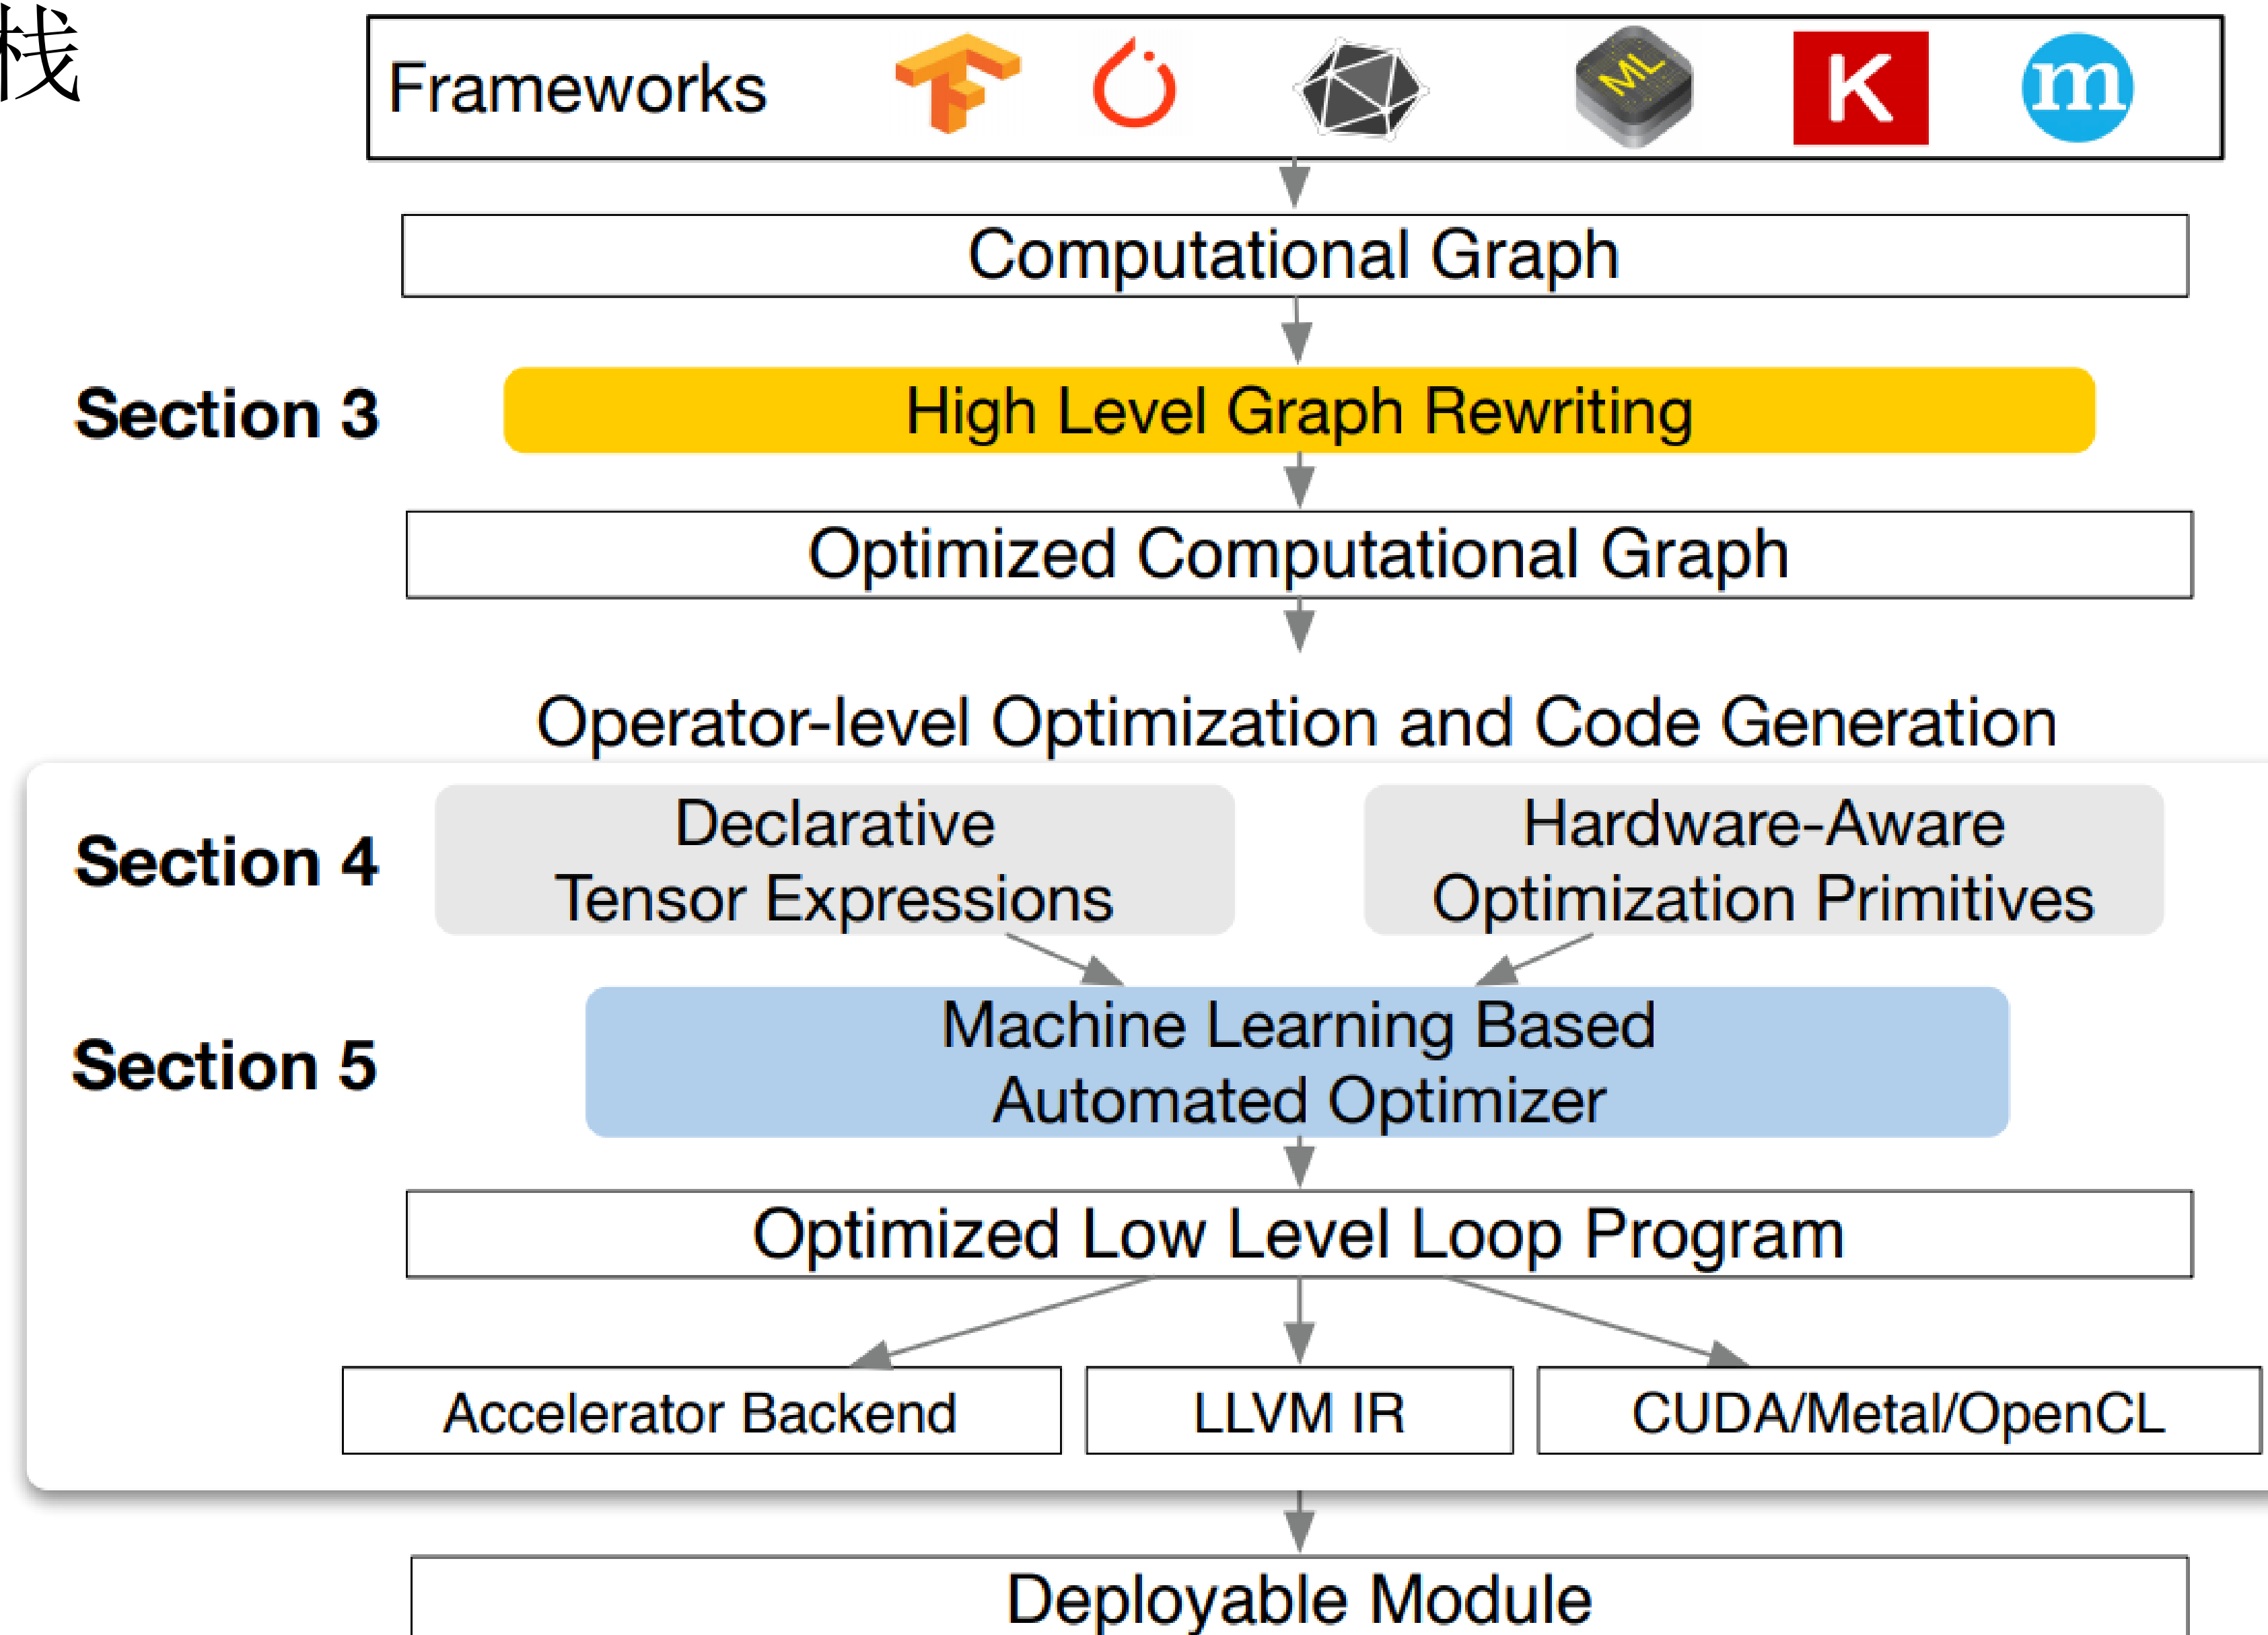
```
mxnet-job-scheduler-0    1/1      Running    0          23s
mxnet-job-server-0       1/1      Running    0          23s
mxnet-job-worker-0       1/1      Running    0          23s
leisu@desk-73:/extend/Workspace/myproj/src/github.com/kubeflow/mxnet-operator/examples/mxnet-operator.v2/train$ kubectl logs mxne
-0
INFO:root:start with arguments Namespace(add_stn=False, batch_size=64, disp_batches=10, dtype='float32', gc_threshold=0.5, gc_tyr
us='0', initializer='default', kv_store='dist_device_sync', load_epoch=None, loss='', lr=0.05, lr_factor=0.1, lr_step_epochs='10'
_size=0, model_prefix=None, mom=0.9, monitor=0, network='mlp', num_classes=10, num_epochs=10, num_examples=60000, num_layers=2, c
d', test_io=0, top_k=0, warmup_epochs=5, warmup_strategy='linear', wd=0.0001)
DEBUG:urllib3.connectionpool:Starting new HTTP connection (1): yann.lecun.com
DEBUG:urllib3.connectionpool:http://yann.lecun.com:80 "GET /exdb/mnist/train-labels-idx1-ubyte.gz HTTP/1.1" 200 28881
DEBUG:urllib3.connectionpool:Starting new HTTP connection (1): yann.lecun.com
DEBUG:urllib3.connectionpool:http://yann.lecun.com:80 "GET /exdb/mnist/train-images-idx3-ubyte.gz HTTP/1.1" 200 9912422
DEBUG:urllib3.connectionpool:Starting new HTTP connection (1): yann.lecun.com
DEBUG:urllib3.connectionpool:http://yann.lecun.com:80 "GET /exdb/mnist/t10k-labels-idx1-ubyte.gz HTTP/1.1" 200 4542
DEBUG:urllib3.connectionpool:Starting new HTTP connection (1): yann.lecun.com
DEBUG:urllib3.connectionpool:http://yann.lecun.com:80 "GET /exdb/mnist/t10k-images-idx3-ubyte.gz HTTP/1.1" 200 1648877
INFO:root:Epoch[0] Batch [10]      Speed: 31347.93 samples/sec      accuracy=0.332386
INFO:root:Epoch[0] Batch [20]      Speed: 41367.77 samples/sec      accuracy=0.709375
INFO:root:Epoch[0] Batch [30]      Speed: 42556.11 samples/sec      accuracy=0.775000
INFO:root:Epoch[0] Batch [40]      Speed: 43240.25 samples/sec      accuracy=0.831250
INFO:root:Epoch[0] Batch [50]      Speed: 37664.58 samples/sec      accuracy=0.873437
INFO:root:Epoch[0] Batch [60]      Speed: 39177.36 samples/sec      accuracy=0.853125
INFO:root:Epoch[0] Batch [70]      Speed: 36217.31 samples/sec      accuracy=0.860938
INFO:root:Epoch[0] Batch [80]      Speed: 36088.77 samples/sec      accuracy=0.865625
INFO:root:Epoch[0] Batch [90]      Speed: 35676.75 samples/sec      accuracy=0.884375
INFO:root:Epoch[0] Batch [100]     Speed: 39835.49 samples/sec      accuracy=0.909375
INFO:root:Epoch[0] Batch [110]     Speed: 38521.27 samples/sec      accuracy=0.892188
INFO:root:Epoch[0] Batch [120]     Speed: 36256.44 samples/sec      accuracy=0.893750
INFO:root:Epoch[0] Batch [130]     Speed: 36556.65 samples/sec      accuracy=0.906250
INFO:root:Epoch[0] Batch [140]     Speed: 34758.38 samples/sec      accuracy=0.921875
INFO:root:Epoch[0] Batch [150]     Speed: 39165.36 samples/sec      accuracy=0.917188
INFO:root:Epoch[0] Batch [160]     Speed: 38720.19 samples/sec      accuracy=0.907813
INFO:root:Epoch[0] Batch [170]     Speed: 36834.55 samples/sec      accuracy=0.900000
INFO:root:Epoch[0] Batch [180]     Speed: 37426.69 samples/sec      accuracy=0.917188
INFO:root:Epoch[0] Batch [190]     Speed: 36872.50 samples/sec      accuracy=0.925000
INFO:root:Epoch[0] Batch [200]     Speed: 40973.13 samples/sec      accuracy=0.914062
INFO:root:Epoch[0] Batch [210]     Speed: 42677.90 samples/sec      accuracy=0.907813
INFO:root:Epoch[0] Batch [220]     Speed: 38584.39 samples/sec      accuracy=0.917188
```



# 什么是TVM

- 一个端到端的深度学习编译软件栈
  - 为CPU, GPU和其他硬件加速器优化深度学习计算负载
  - 自动转换计算图以最小化内存使用, 优化数据布局, 算子融合
  - 提供从现有的前端框架到裸机硬件的端到端编译
- 官网地址:
  - <https://tvm.ai/>

# TVM软件栈





# TVM使用

## Typical Workflow of NNVM Compiler

**model from  
framework**

```
graph, params = nnvm.frontend.from_xyz(...)
```

**compile**

```
graph, lib, params = nnvm.compiler.build(  
    graph, target="cuda", {"data", data_shape}, params=params)
```

**deploy**

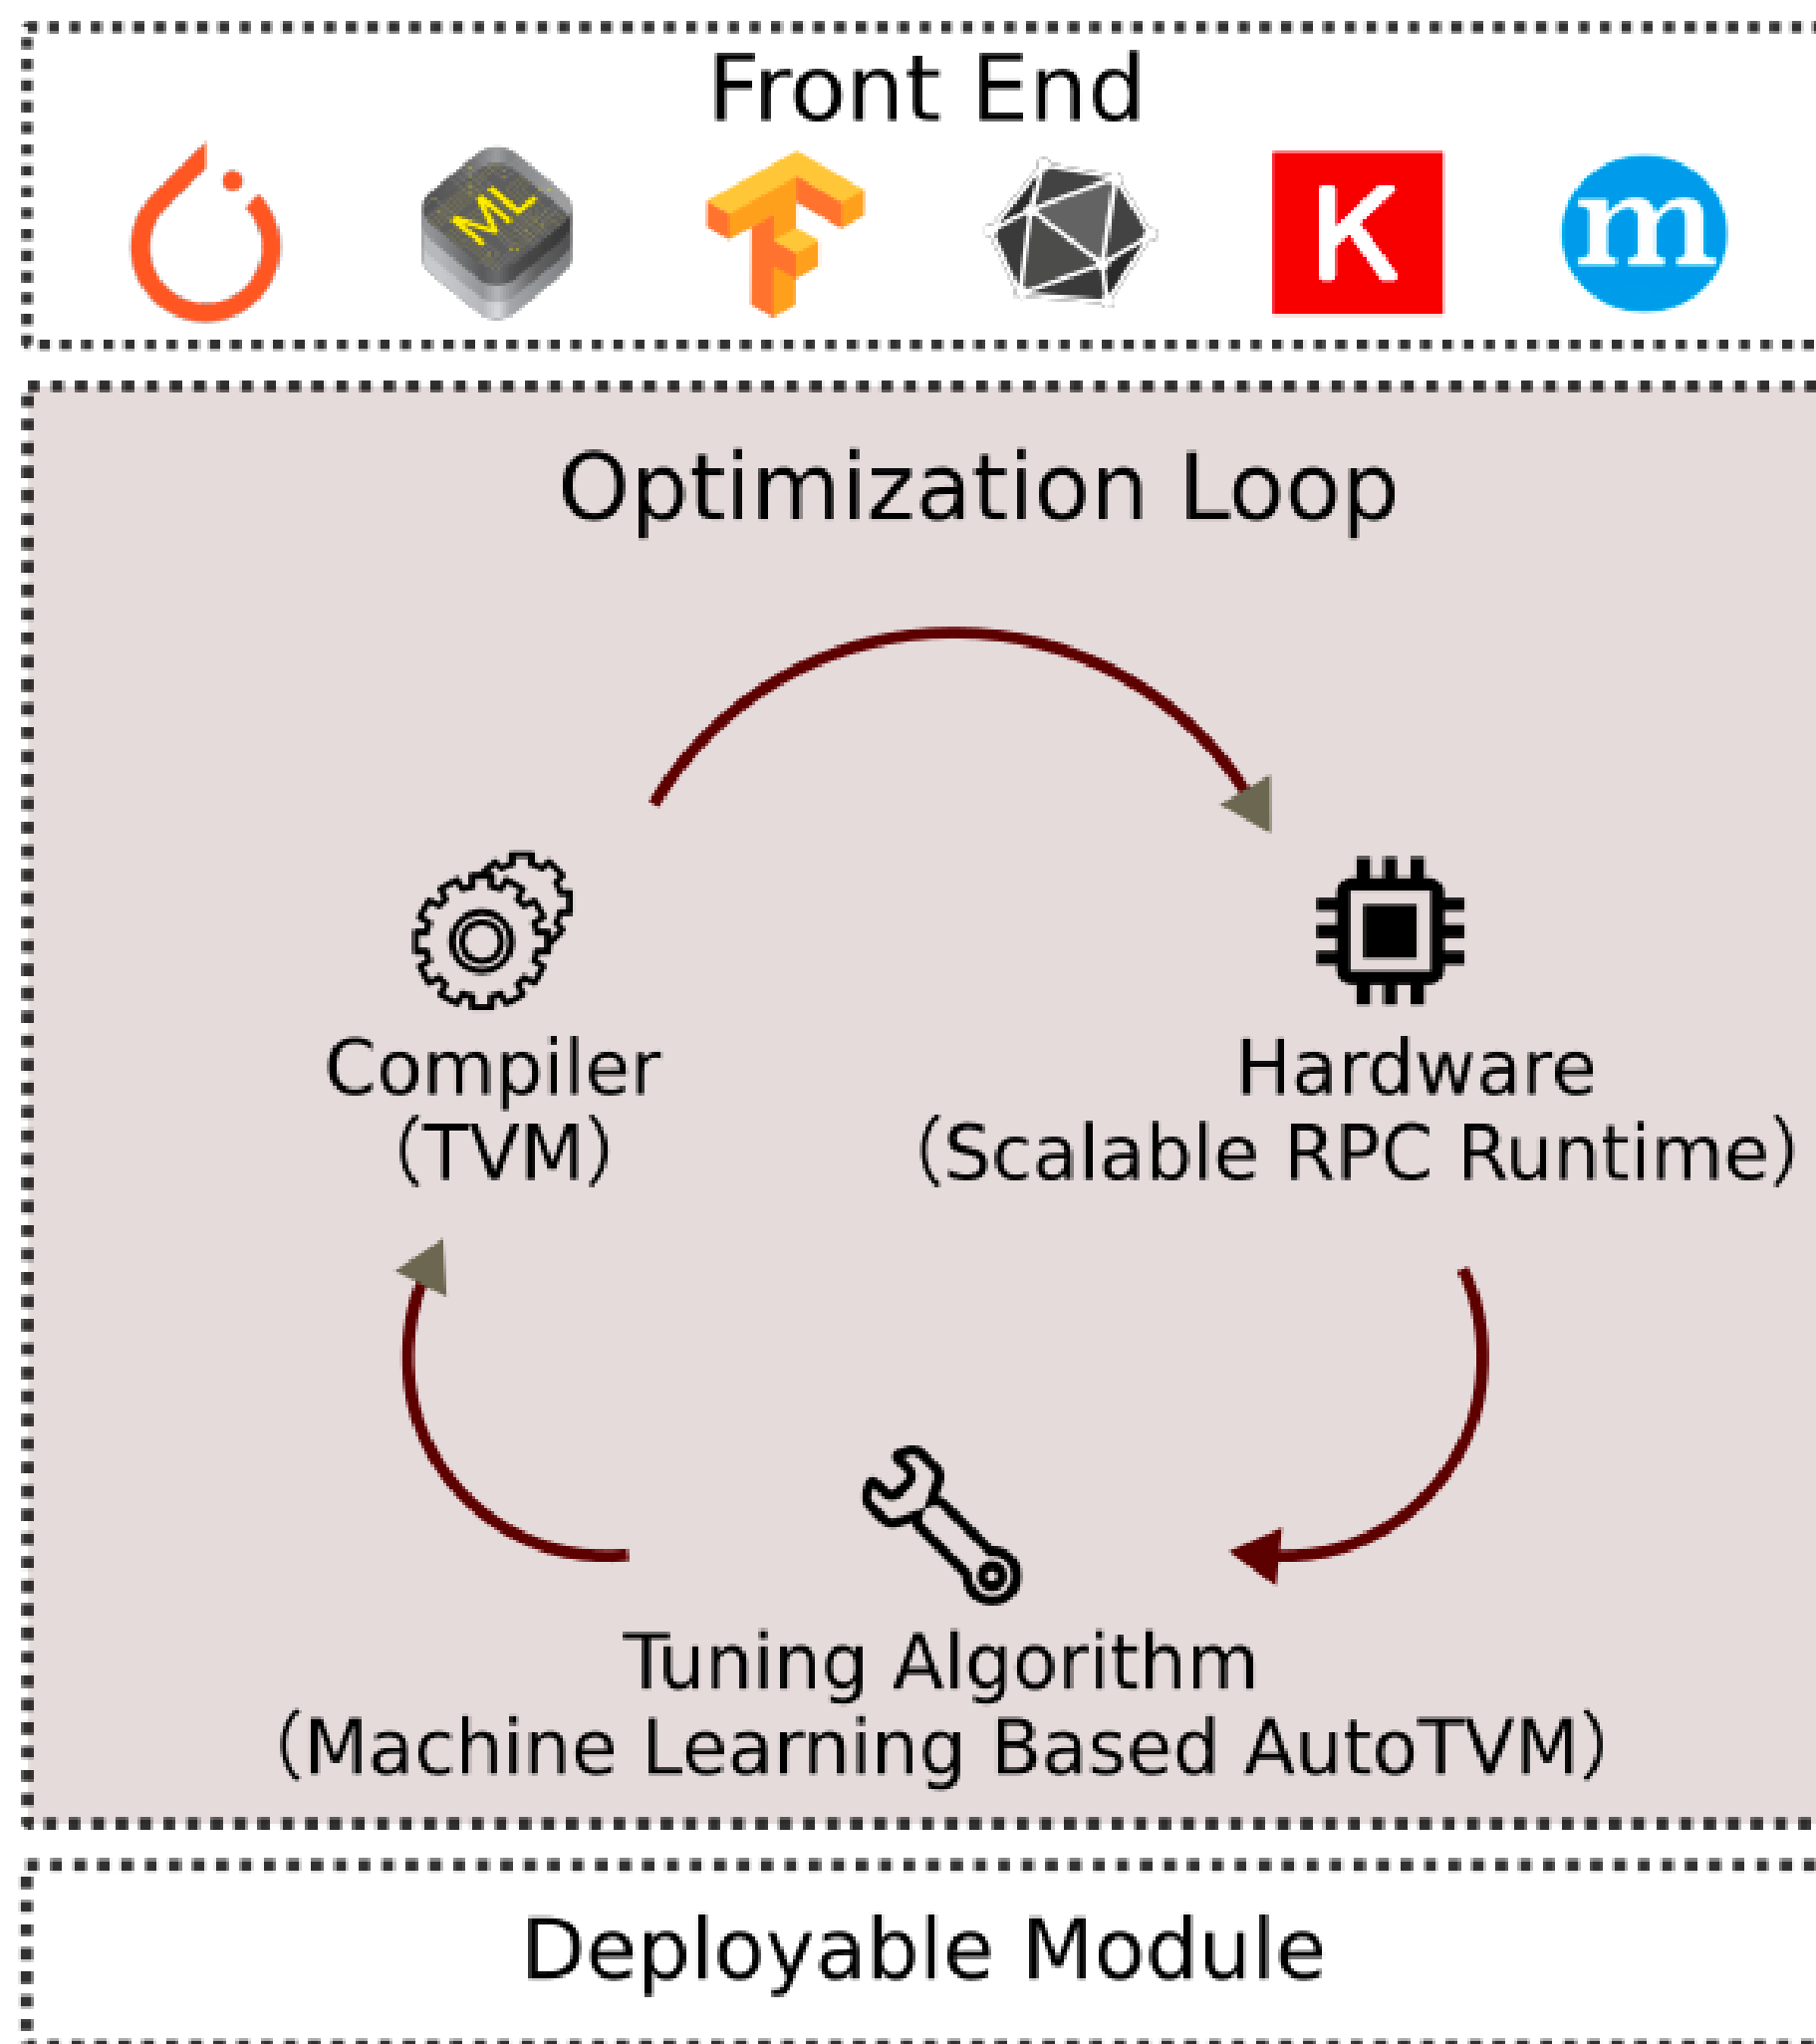
```
module = graph_runtime.create(graph, lib, tvn.gpu(0))  
module.set_input(**params)  
module.run(data=data_array)  
output = tvn.nd.empty(out_shape, ctx=tvn.gpu(0))  
module.get_output(0, output)
```

# 什么是AutoTVM

- 痛点：
  - 高效的operator受多种因素影响：硬件设备的类型、输入大小、tensor layout等，手工优化费时费力。
  - 参数搜索空间巨大
- 方案：
  - 采用机器学习方法自动寻找较优的参数配置





# AutoTVM机制



- 基本流程
  - 从前端框架获取和生成所有操作符的kernel
  - Tuner从可调参数搜索空间里寻找并编译出一批有提速潜质的kernel，通过RPC调用在硬件上执行这些kernel
  - Tuner获得相应kernel的性能数据并作为训练数据训练新的预测模型
  - 获得新的预测模型后，Tuner根据新的预测结果寻找下一批有提速潜质的kernel实现，并循环上述过程
- 组件：
  - Tunertracker：将任务调度至Tunerserver
  - Tunerserver：在实际硬件上执行kernel并返回性能数据
  - Tuner：实际的tuning程序，基于机器学习算法寻找较优配置并生成kernel

# AutoTVM效果



# 使用mxnet-operator提交auto-tuning任务

```
apiVersion:
"kubeflow.org/v1alpha2"
kind: "MXJob"
metadata:
  name: "auto-tuning-job"
spec:
  jobMode: MXTune
  mxReplicaSpecs:
    TunerTracker:
      replicas: 1
      restartPolicy: Never
      template:
        spec:
          containers:
            - name: mxnet
              image: mxjob/auto-
tuning:gpu
              command: ["python3"]
              args:
["/home/scripts/start-job.py"]
          resources:
            limits:
nvidia.com/gpu: 1

TunerServer:
  label: 2080ti
  replicas: 1
  restartPolicy: Never
  template:
    spec:
      containers:
        - name: mxnet
          image:
mxjob/auto-tuning:gpu
          command:
["python3"]
          args:
["/home/scripts/start-
job.py"]
          resources:
            limits:
nvidia.com/gpu: 1

Tuner:
  replicas: 1
  restartPolicy: Never
  template:
    spec:
      containers:
        - name: mxnet
          image:
mxjob/auto-tuning:gpu
          command:
["python3"]
          args:
["/home/scripts/start-
job.py"]
```



# 使用mxnet-operator提交auto-tuning任务

```
if __name__ == '__main__':
    mx_config = json.loads(os.environ.get('MX_CONFIG') or '{}')
    cluster_config = mx_config.get('cluster', {})
    labels_config = mx_config.get('labels', {})
    task_config = mx_config.get('task', {})
    task_type = task_config.get('type')
    task_index = task_config.get('index')

    if task_type == "":
        print("No task_type, Error")
    elif task_type == "tunertracker":
        addr = cluster_config["tunertracker"][0]
        command = "python3 -m tvn.exec.rpc_tracker --port={0}".format(addr.get('port'))
        print("DO: " + command)
        os.system(command)
    elif task_type == "tunerserver":
        time.sleep(5)
        addr = cluster_config["tunertracker"][0]
        label = labels_config["tunerserver"]
        command = "python3 -m tvn.exec.rpc_server --tracker={0}:{1} --key={2}".format(addr.get('url'), addr.get('port'), label)
        print("DO: " + command)
        os.system(command)
    elif task_type == "tuner":
        time.sleep(5)
        addr = cluster_config["tunertracker"][0]
        label = labels_config["tunerserver"]
        command = "python3 /home/scripts/auto-tuning.py --tracker {0} --tracker_port {1} --server_key {2}".format(addr.get('url'),
addr.get('port'), label)
        print("DO: " + command)
        os.system(command)
    else:
        print("Unknow task type! Error")
```

# 使用mxnet-operator提交auto-tuning任务

```
leisu@desk-73:/extend/Workspace/myproj/src/github.com/kubeflow/mxnet-operator/examples/mxnet-operator.v2/tune$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
auto-tuning-job-tuner-0	1/1	Running	0	9m
auto-tuning-job-tunerserver-0	1/1	Running	0	9m
auto-tuning-job-tunertracker-0	1/1	Running	0	9m

```
leisu@desk-73:/extend/Workspace/myproj/src/github.com/kubeflow/mxnet-operator/examples/mxnet-operator.v2/tune$ kubectl logs auto-tuning-job-tuner-0
```

Extract tasks...

...100%, 0.02 MB, 23 KB/s, 1 seconds passed

Tuning...

[Task 1/12] Current/Best: 2343.76/2632.19 GFLOPS | Progress: (100/100) | 273.24 s Done.

[Task 2/12] Current/Best: 285.03/ 392.79 GFLOPS | Progress: (100/100) | 123.07 s Done.

# Mxnet-operator Roadmap

- 提高AutoTVM运行效率
- MXJob支持MXBoard
- 集成Kube-batch



# 致谢



姜宗沛  
清华大学



杨瑞恒  
上海交通大学



# Q & A



扫码关注公众号,获取 316 北京站 PPT